

Lab package 3: Supervised classification and regression

Wednesday, October 6, 2021 1:43 PM

Supervised classification is in general easier to understand than unsupervised classification. Data for unsupervised classification have measures for features (X). Data for supervised classification have measures for features (X) and the true labels (y).

The general procedures for supervised classification:

1. Clean data
2. Select appropriate features as the input variables
3. Randomly split data into a training set and a test set: commonly 70/30, 75/25, or 80/20 split
4. Scale and transform data:
 - a. Use training features to fit a scaler
 - b. Transform both training features and test features.
5. Select, initial, and fit a machine learning model with training features and training labels
6. Fit the model with test features to predict test labels
7. Compare test labels and predict test labels to evaluate the model performance
8. Calculate additional evaluation metrics
9. Run different models to select the one with the best performance
10. Some may opt to use the entire data set to finalize the model (especially for software production)

Strictly speaking, the training procedures should include two separate data sets: training data and validation data. A more rigorous approach includes the following:

1. Further split the training data into training data and validation data (commonly, 80/20). If considered validation, we usually have the training/testing split to 80/20; and then have training/validation split to 80/20; resulting training/validation/testing split at 60/20/20
2. Most machine learning methods involve training hyperparameters and parameters. Validation data are used as a part of the training process to fine tune parameters and assess the model skills.

Classification is to classify observations into discrete groups (y) based on measurable features (X). Groups are discrete labels. Regression is to make predictions (y) based on measurable features (X). Predictions are continuous values.

Procedures wise, classification and regression are similar, but mathematically they are distinct.

We will practice supervised classification and regression in Lab package 3.

Objectives:

1. Are there some census tracts with distinctively more population of high cholesterol than the other tracts in the city of Dallas?
2. If so, how well can we group the census tracts according to the percentage of people with high cholesterol in each tract, and what is the optimal number of groups?
3. How may food consumptions predict which group of percent population with high cholesterol for a census tract in Dallas?
4. How may food consumptions predict the percent population with high cholesterol in Dallas?

The python script:

```
[1] import pandas as pd

[2] !pip install sodapy
from sodapy import Socrata
client = Socrata("chronicdata.cdc.gov", None)
results=client.get("cwsq-ngmh", where="stateabbr = 'TX'", limit=200000)
df = pd.DataFrame.from_records(results)

[3] df

[4] counties = ['Dallas', 'Collin', 'Denton', 'Rockwell', 'Kaufman']
dfcounties = df[df['countyname'].isin(counties)]
dfcounties

[5] !pip install geopandas

[6] import geopandas as gpd
from shapely.geometry import shape

[7] dfcounties['geolocation'].dtype

dtype('O')

[8] dfcounties['geom'] = dfcounties['geolocation'].apply(shape)
gdf = gpd.GeoDataFrame(dfcounties).set_geometry('geom')

[9] gdf.crs

[10] file = r'//content/drive/MyDrive/package1/DallasTract2020.shp.zip'
DallasTracts = gpd.read_file(file)

[11] DallasTracts.crs

[12] DallasTracts = DallasTracts.to_crs(epsg=2276)
```

Q#1: Which of the following statement is false:

- A. Clustering analysis is an unsupervised classification
- B. Hotspot analysis is an unsupervised classification
- C. Facial recognition is a supervised classification
- D. Toponym resolution is a supervised classification
- E. None of the above

Q#2: Which of the following statement is false in comparison between training and validation during machine learning:

- A. The training explores the space of a chosen hypothesis class
- B. The validation seeks the best possible hypothesis in a chosen hypothesis class
- C. More data samples should be used for training than for validation
- D. Validation data are part of the training data
- E. None of the above

Q#3: If we want to do a 5-fold cross-validation, what is the proportion of training data to validation data?

- A. 5 to 1
- B. 4 to 1
- C. 3 to 2
- D. 2 to 3
- E. All of the above are possible

Q#4: The health data are from the Center for Disease Control <https://dev.socrata.com/foundry/chronicdata.cdc.gov/cwsq-ngmh>. The data include _____ (a number) measures at four geographic area units: county, place, census tracts, and _____. The measures are calculated from three sources. Specifically, high cholesterol estimates are from _____ (a year) _____ (a data source - using the acronym).

Q#5: Why do we need to do apply(shape) in line [8]?

- A. Assign geolocation to a shape
- B. Convert geolocation from data type object to data type shape
- C. Create a new data type shape from geolocation
- D. Transform the shape of geolocation
- E. None of the above

Q#6: DallasHealth has _____ (a number) records and _____ (a number) features. There are _____ (a number) of records with missing data. The field _____ (a field name) indicates the records with data values for high cholesterol.

```

[13] gdf = gdf.set_crs(eps=4326)
gdf = gdf.to_crs(eps=2276)

[14] !apt install libspatialindex-dev
!pip install rtree

[15] gdf.crs

[16] DallasHealth = gpd.clip(gdf, DallasTracts)

[17] base = DallasTracts.plot()
DallasHealth.plot(ax=base, marker='o', color='red', markersize=5)

[18] columns=['locationname', 'category', 'measure','data_value', 'low_confidence_limit','high_confidence_limit',
'totalpopulation', 'locationid', 'categoryid','short_question_text','geom']

[19] DallasHealth = DallasHealth[columns]

[20] DallasHealth.isnull().sum()

[21] %load_ext google.colab.data_table

[22] DallasHighChol = DallasHealth[DallasHealth['short_question_text'] == 'High Cholesterol']

[23] !pip install libpysal

[24] DallasHealth.dtypes

[25] DallasTracts.columns

[26] DallasTractsOnly = DallasTracts[['id', 'name', 'geometry']]
TractHighChol = pd.merge(DallasTractsOnly, DallasHighChol, how='left', left_on='id', right_on='locationname')
TractHighChol.plot(column='data_value', cmap='OrRd', edgecolor='k')

[27] !pip install mapclassify

[28] TractHighChol.isnull().sum()

[29] TractHighChol.dtypes

[30] TractHighChol['data_value'] = pd.to_numeric(TractHighChol['data_value'], downcast='float')

[31] missing_kwds={'color':'lightgrey', 'edgecolor':'k', 'hatch': '/', 'edgecolor':'red', 'label': 'Missing values'}
legend_kwds = {'loc': 'lower right'}
TractHighChol.plot(column='data_value', scheme='quantiles', figsize=(15, 10), missing_kwds=missing_kwds, legend=True, legend_kwds=legend_kwds)

[32] TractHighChol.plot(column='data_value', scheme='fisherjenks', k=2, figsize=(15, 10), missing_kwds=missing_kwds, legend=True, legend_kwds=legend_kwds)

[33] TractHighChol['class'] = [1 if x > 31.30 else 0 for x in TractHighChol['data_value']]

[34] TractHighChol.groupby(['class']).size()

[35] file = r'./content/drive/MyDrive/package2/DallasConsumptions.xlsx'
DallasConsump = pd.read_excel(file)
DallasConsump

[36] DallasConsump.columns

[37] cols = ['GEOID', '2021MedianHouseholdIncome', '2021AlcoholicBevs', '2021LunchatVendMobileVendors', '2021LunchatFastFoodTake-OutDeliv',
'2021DinneratFastFoodTake-OutDeliv', '2021DinneratVendMobileVendors', '2021MeatPoultryFishEggs', '2021FreshVegetables', '2021FreshFruit']

```

Q#7: Why do pd.to_numeric in line [30]?

- A. We need to downcast data_value from float to integer.
- B. We need to downcast data_value from integer to float.
- C. We need to downcast data_value from string to float.
- D. We need to downcast data_value from object to float.

```

[38] DallasConsump = DallasConsump[cols]

[39] TractHighChol['id'] = pd.to_numeric(TractHighChol['id'])

[40] TractHighCholConsump = pd.merge(TractHighChol, DallasConsump, how='left', left_on='id', right_on='GEOID')

[41] TractHighCholConsump.columns

[42] TractHighCholConsump.isnull().sum()

[43] TractHighCholConsump.dropna(how='any', inplace=True)

import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from imblearn.over_sampling import RandomOverSampler
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import confusion_matrix
from sklearn import metrics

[45] X = TractHighCholConsump[['2021MedianHouseholdIncome',
    '2021AlcoholicBevs', '2021LunchatVendMobileVendors',
    '2021LunchatFastFoodTake-OutDeliv', '2021DinneratFastFoodTake-OutDeliv',
    '2021DinneratVendMobileVendors', '2021MeatPoultryFishEggs',
    '2021FreshVegetables', '2021FreshFruit']]
    y = TractHighCholConsump['class']

[46] X.shape

```

```

[47] Xtrain, Xtest, ytrain, ytest = train_test_split(X, y, test_size=0.2, random_state=1)

[48] scaler = StandardScaler().fit(Xtrain)
    XsTrain = scaler.transform(Xtrain)
    XsTest = scaler.transform(Xtest)

[49] resampler = RandomOverSampler()
    XsTrain_over, ytrain_over = resampler.fit_resample(XsTrain, ytrain)

np.unique(ytrain_over, return_counts=True)

[51] lrclf = LogisticRegression(random_state=10).fit(XsTrain, ytrain)
    ylrpred = lrclf.predict(XsTest)
    confmat = confusion_matrix(ytest, ylrpred)
    confmat

[52] tn, fp, fn, tp = confusion_matrix(ytest, ylrpred).ravel()
    print('tn:', tn)
    print('fp:', fp)
    print('fn:', fn)
    print('tp:', tp)

[53] print('accuracy:', metrics.accuracy_score(ytest, ylrpred))
    print("recall:", metrics.recall_score(ytest, ylrpred))
    print("precision:", metrics.precision_score(ytest, ylrpred))
    print('F-score:', metrics.f1_score(ytest, ylrpred))

```

Q#8: After train_test_split, we have _____ (a number) training samples and _____ (a number) test samples. There are _____ (more or less) training samples in class 0 than in class 1. Based on the numbers of training samples in classes 0 and 1, if an algorithm blindly makes all predictions to class 1, the algorithm's prediction accuracy will be _____ (a ratio, round up to two digits, such as 90%).

- Q#9: Which of the statements is false:
- A. If training data have high imbalanced distributions among classes, the prediction results are easier to get high accuracy.
 - B. If a prediction has high accuracy, we can have high trust in the prediction skills of the algorithm.
 - C. A prediction can be at a high accuracy but low precision if the prediction produces high false positive cases.
 - D. A prediction has a high recall if the algorithm is good at predicting positive cases.
 - E. None of the above.

```
[54] def clfsummary(test, pred):
    tn, fp, fn, tp = confusion_matrix(test, pred).ravel()
    print('tn:', tn)
    print('fp:', fp)
    print('fn:', fn)
    print('tp:', tp)
    print('accuracy:', metrics.accuracy_score(test, pred))
    print("recall:", metrics.recall_score(test, pred))
    print("precision:", metrics.precision_score(test, pred))
    print('F-score:', metrics.f1_score(test, pred))

[55] svmclf = SVC(random_state=20, kernel='linear', verbose=True).fit(XsTrain_over, ytrain_over)
ysvpred = svmclf.predict(XsTest)
confmat = confusion_matrix(ytest, ysvpred)
confmat

[56] clfsummary(ytest, ysvpred)

[57] rfclf = RandomForestClassifier(n_estimators = 5, max_depth=3, random_state=10).fit(XsTrain_over, ytrain_over)
yrfpred = rfclf.predict(XsTest)
confmat = confusion_matrix(ytest, yrfpred)
confmat

[58] clfsummary(ytest, yrfpred)

[59] nnclf = MLPClassifier(random_state=30, hidden_layer_sizes=(5, 3)).fit(XsTrain_over, ytrain_over)
ynnpred = nnclf.predict(XsTest)
confmat=confusion_matrix(ytest, ynnpred)
confmat

[60] clfsummary(ytest, ynnpred)

[61] from sklearn.model_selection import GridSearchCV
```

```
[62] lrcf = LogisticRegression(random_state=10)
param_grid = dict(C=[-1, 0, 1, 5, 10], solver=['newton-cg', 'sag', 'saga', 'lbfgs'])
grid = GridSearchCV(lrcf, param_grid, cv=5, scoring='accuracy', n_jobs=-1, return_train_score=True)
grid.fit(XsTrain_over, ytrain_over)
GridSearchCVout = pd.DataFrame(grid.cv_results_)[['mean_test_score', 'std_test_score', 'params']]
GridSearchCVout

[63] lrcf = LogisticRegression(C=10, solver='lbfgs').fit(XsTrain_over, ytrain_over)
ylr2pred = lrcf.predict(XsTest)
confmat = confusion_matrix(ytest, ylr2pred)
confmat

[64] clfsummary(ytest, ylr2pred)

[65] from sklearn.neural_network import MLPRegressor

[66] y = TractHighCholConsump['data_value']
Xtrain, Xtest, ytrain, ytest = train_test_split(X, y, test_size=0.2, random_state=5)
scaler = StandardScaler().fit(Xtrain)
XsTrain = scaler.transform(Xtrain)
XsTest = scaler.transform(Xtest)
resampler = RandomOverSampler()
XsTrain_over, ytrain_over = resampler.fit_resample(XsTrain, ytrain)

[67] mlpreg = MLPRegressor(random_state=40)
param_grid = dict(hidden_layer_sizes=[(2),(3), (4,2), (5,3)],
                    activation=['logistic', 'relu', 'tanh'],
                    solver=['lbfgs', 'sgd', 'adams'])
grid = GridSearchCV(mlpreg, param_grid, cv=5, scoring='max_error', n_jobs=-1, return_train_score=True)
grid.fit(XsTrain_over, ytrain_over)
GridSearchCVout = pd.DataFrame(grid.cv_results_)[['mean_test_score', 'std_test_score', 'params']]
GridSearchCVout
```

```
[68] mlpreg = MLPRegressor(random_state=40, activation='logistic', hidden_layer_sizes=(5,3), solver='sgd', max_iter=50000).fit(XsTrain_over, ytrain_over)
ymlpregpred = mlpreg.predict(XsTest)
mae=metrics.mean_absolute_error(ytest, ymlpregpred)
mse = metrics.mean_squared_error(ytest, ymlpregpred)
r2 = metrics.r2_score(ytest, ymlpregpred)
explained_variance = metrics.explained_variance_score(ytest, ymlpregpred, multioutput='uniform_average')

print("The model performance for testing set")
print("-----")
print('Mean Absolute Error is {}'.format(mae))
print('Mean Squared Error is {}'.format(mse))
print('R2 score is {}'.format(r2))
print('Explained Variance Score is {}'.format(explained_variance))

[69] from sklearn.linear_model import LinearRegression

[70] linreg = LinearRegression().fit(XsTrain_over, ytrain_over)
ylinregpred = linreg.predict(XsTest)
mae=metrics.mean_absolute_error(ytest, ylinregpred)
mse = metrics.mean_squared_error(ytest, ylinregpred)
r2 = metrics.r2_score(ytest, ylinregpred)
explained_variance = metrics.explained_variance_score(ytest, ylinregpred)

print("The model performance for testing set")
```

Q#10: Which of the statements is false based on the MLPregressor and the linear regressor in lines 68-70:

- A. The MLPregressor predicts better than the Linear regressor.
- B. The MLPregressor consists of many logistic regressors.
- C. The Linear regressor includes 9 independent variables.
- D. The Linear regressor can only explain about 20% of the variance in the data.

```
mse = metrics.mean_squared_error(ytest, ylinregpred)
r2 = metrics.r2_score(ytest, ylinregpred)
explained_variance = metrics.explained_variance_score(ytest, ylinregpred)

print("The model performance for testing set")
print("-----")
print('Mean Absolute Error is {}'.format(mae))
print('Mean Squared Error is {}'.format(mse))
print('R2 score is {}'.format(r2))
print('Explained Variance Score is {}'.format(explained_variance))
```

- B. The MLPregressor consists of many logistic regressors.
- C. The Linear regressor includes 9 independent variables.
- D. The Linear regressor can only explain about 20% of the variance in the test data.
- E. None of the above.