

MOTIVATION FOR HAMILTON CHAPTER 1

- Data analysis fundamental: Treat your data with respect to learn something about the underlying data generating process.
 - Data tell a story about the phenomena under investigation.
 - Always handle data and analysis results with a critical attitude and use common sense. Link the results back to your original observations.
 - Always ask yourself: Do the data or the generated analysis results make sense?
- **Describing the variability** and distribution of a variable is the **required** first step of any data analysis.
- The **shape** of an univariate distribution can have **substantial impact** on the outcome of statistical procedures.
E.g.: **Outliers** or **heavy tails** may detrimentally influence the outcome of model calibrations and parameter estimations.
- Not accounting for the distribution of variables can force a researcher to redo their data analysis at a later state.
- Most methods assume **symmetric** or preferably **normally** distributed variables.
- Transformations to symmetry are discussed in Chapter 1. Note, statisticians use many more transformations under particular circumstances.
E.g., we will encounter later the logit-transformation.

Central Limit Theorem

- Def. Central Limit Theorem: Let X_1, X_2, \dots, X_n be a **random independent** sample of size n drawn from an **arbitrarily distributed** population with expectation μ and standard deviation σ .


Then for large enough sample sizes n , the sampling distribution of the arithmetic mean \bar{X} is [a] asymptotically (i.e., as the sample size $n \rightarrow \infty$) normal distributed [b] with

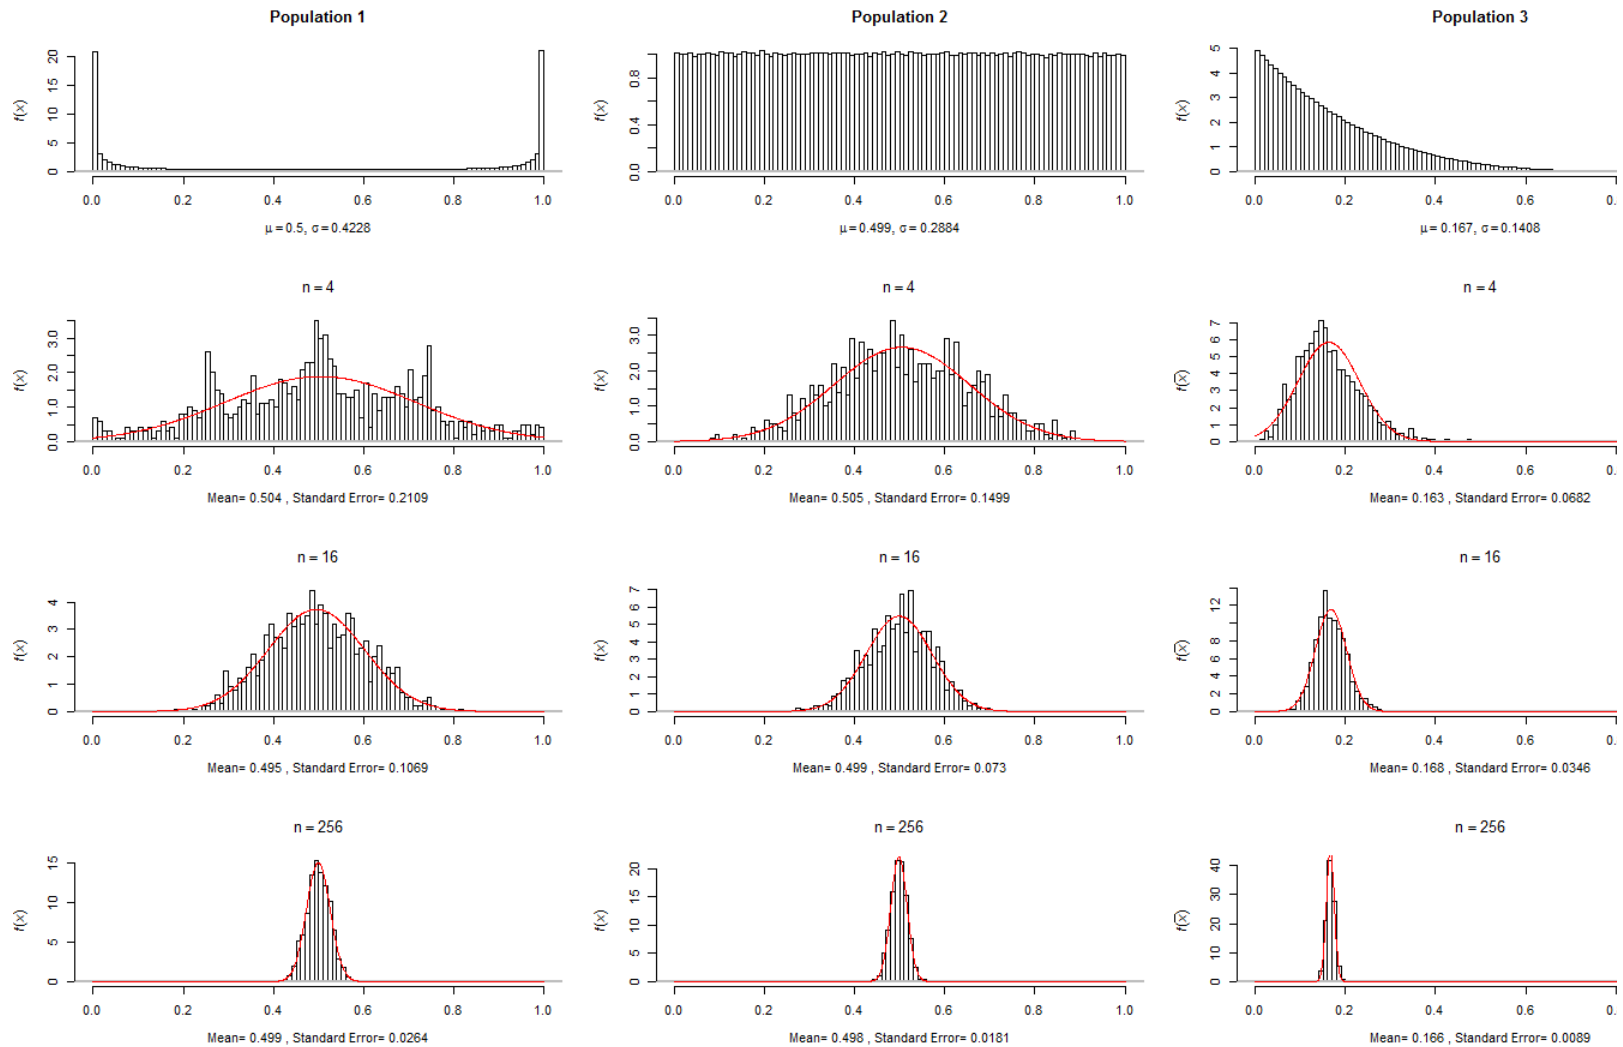
$$\bar{X} \sim N\left(\mu, \frac{\sigma^2}{n}\right).$$

Proof for independent sample observations X_i :

$$\text{Var}\left(\frac{1}{n} \cdot \sum_{i=1}^n X_i\right) = \frac{1}{n^2} \cdot \sum_{i=1}^n \underbrace{\text{Var}(X_i)}_{=\sigma^2} = \frac{1}{n^2} \cdot n \cdot \sigma^2 = \frac{\sigma^2}{n}$$

- Implications:
 - Thus the **standard error** $s_{\bar{X}} = \sigma/\sqrt{n}$ of the mean \bar{X} **shrinks** by $1/\sqrt{n}$ with increasing sample size.
 - Another implication of the central limit theorem is that the **sum of a set of small random errors** or shocks will lead to **normal distributed total error**.
 - In contrast, the product of a set of small random errors will lead to a log-normal distributed total error.

- Example: Central limit theorem with the -script **CENTRALLIMIT.R**:



Review: The Shape of Distributions

- Distributions can be distinguished with regards the **balance** of their left and right tails:
 - **Symmetric** distributions. Tails are balanced into either direction from a central value.
 - **Negatively** skewed distributions (long tail into the negative direction)
 - **Positively** skewed distributions (long tail into the positive direction). These distributions frequently emerge for variables with a **binding lower origin** (like zero income).
 - Extreme skewness may hint at **outliers** that do not match the rest of the observed data.
- The number of meaningful clusters of observations is described by the term modality:
 - Uni-modality refers to just one peak.
 - Bi-modality refers to two outstanding peaks
 - Multimodality refers to more than two outstanding peaks.
- Multimodality may hint at a heterogeneous underlying data generating process in which the underlying process for observations in the first mode is different for observation in the second mode.

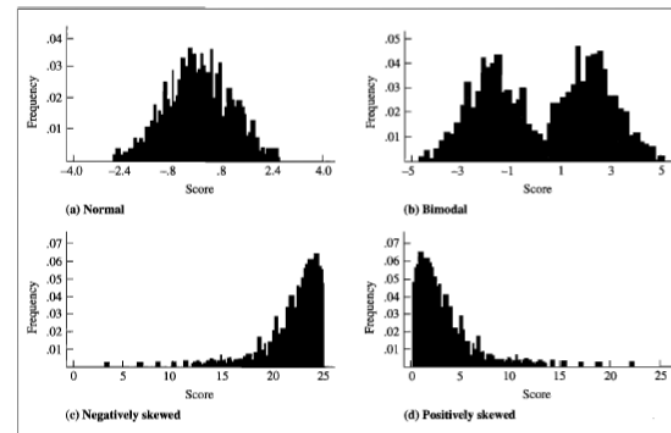


Figure 3.9
Shapes of frequency distributions: (a) Normal; (b) Bimodal; (c) Negatively skewed; (d) Positively skewed

Quantiles and Percentiles

- Technically, quantiles and percentiles are generated from a **sorted list** of the original data points $x_{[1]} \leq x_{[2]} \leq x_{[3]} \leq \dots \leq x_{[n-1]} \leq x_{[n]}$ where each observations has an assigned rank $i \in \{1, 2, \dots, n\}$, with $i = 1$ for the smallest observation and $i = n$ for the largest observation.
- For a give data value $x_{[i]}$ the **percentile** approximates the proportion of sample observations less or equal to $x_{[i]}$, that is, their cumulative distribution:

$$p_{[i]} = \frac{i - \frac{1}{2}}{n} \approx \Pr(X \leq x_{[i]}) = \int_0^{x_{[i]}} f(x) \cdot dx.$$

Note that the $\alpha = 0.5$ of the percentile equation $p(x_{[i]}) = \frac{i - \alpha}{n + (1 - \alpha) - \alpha}$ has been chosen here.

- A **quantile** is the *potentially fictitious* data value of a distribution, which is associated with a particular percentile value.
- Important quantiles are:
 - 0.25 quantile also called Q_1 quartile (25 % of the observations are smaller or equal to this quantile value)
 - 0.50 quantile also called the median (50 % of the observations are smaller or larger than the given quantile value)

- 0.75 quantile also called Q_3 quartile (75 % of the observations are smaller or equal to this quantile value and 25 % of the observations are larger than this value)
- A measure of spread is the inter-quartile range: $IQR = Q_3 - Q_1$


Box-Plots

- Construction of the box-plot
 - Draw a **box** from Q_1 to Q_3 . Mark the **median** Q_2 in the center of the box with a line.
 - Definition of **adjacent values** $x_{low}^{adj} = \min(x_{[i]} \in (Q_1, Q_1 + 1.5 \cdot IQR) \text{ plus } x_{[i]} \text{ in dataset})$ and $x_{high}^{adj} = \max(x_{[i]} \in (Q_3, Q_3 + 1.5 \cdot IQR) \text{ plus } x_{[i]} \text{ in dataset})$.

The term $x \in (a, b)$ means, all x -values in the interval between a and b .

Draw the “fences” so they just include the smallest and largest data values x_{low}^{adj} and x_{high}^{adj} , respectively.

- **Outliers** are in the interval $[1.5 \cdot IQR, 3.0 \cdot IQR]$ starting from Q_1 below or Q_3 above, respectively.
Severe outliers are beyond that range ($> 3.0 \cdot IQR$)



- Use of box-plots:
 - Easy visual description of the distribution of a variable and potential outliers
 - Comparison of distributions for several variables side-by-side.

QUANTILE-NORMAL PLOT

- Calculate the theoretical quantiles of a normally distributed random variable $Y_{[i]}$ (assuming the mean μ and the variance σ^2 were estimated from the sample data) based on the given sample percentiles $p_{[i]}$ of the observed variable $x_{[i]}$.
- **Quantile-Normal Plot:** Plot the theoretical normal distribution quantiles $Y_{[i]}$ on the abscissa (X-axis) against their matching empirical distribution of $x_{[i]}$ on the ordinate (Y-axis).

Interpretation:

- Diagonal with slope 1 => equal distributions.
- Not a straight-line => different shapes.

PROPERTIES OF ARITHMETIC MEAN

- Implications of the **zero-sum** property

$$\sum_{i=1}^n (Y_i - \bar{Y}) = 0$$
 Assuming the mean is known, then

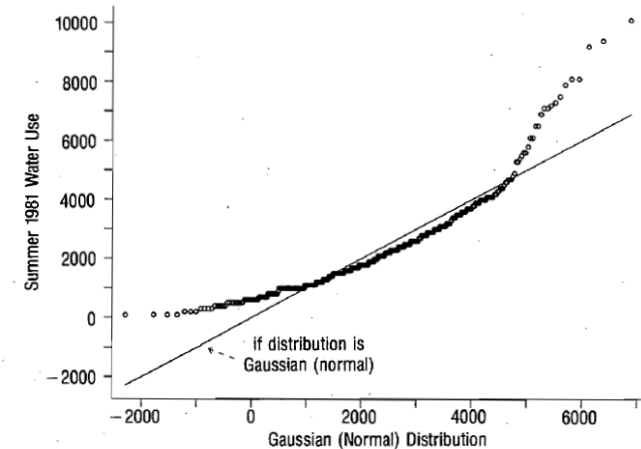


Figure 1.9 Quantile-normal plot of household water use (positively skewed).

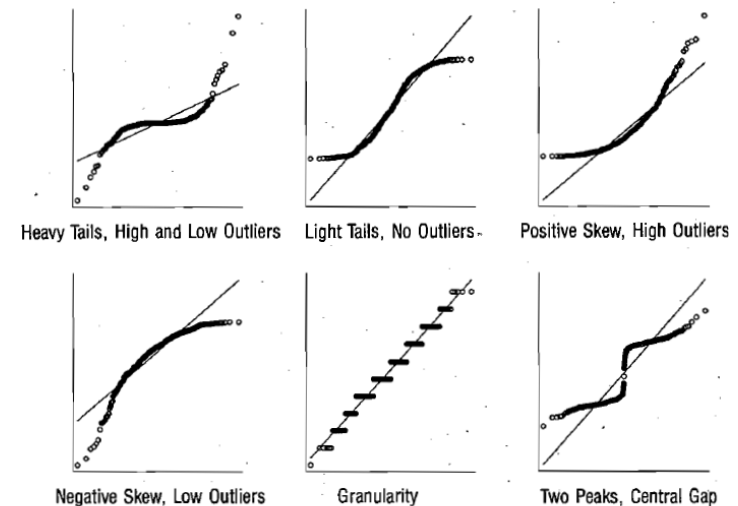


Figure 1.10 Quantile-normal plots reflect distribution shape.

$n - 1$ observation can vary freely, whereas we can predict the last observation with certainty.

$$\begin{aligned}\sum_{i=1}^n (Y_i - \bar{Y}) &= 0 \\ \Rightarrow \sum_{i=1}^n Y_i &= n \cdot \bar{Y} \\ \Rightarrow Y_n &= n \cdot \bar{Y} - \sum_{i=1}^{n-1} Y_i\end{aligned}$$

That implies that we *lose one degree of freedom*.


- Implication of the **least squares property** $\min_{\theta} \sum_{i=1}^n (Y_i - \theta)^2 \Rightarrow \theta = \bar{Y}$.

Large deviations have a strong impact on the estimated mean, variance etc. because the large deviations are squared

\Rightarrow Thus, large deviations pull the mean into their direction.

\Rightarrow Standard deviations are drastically inflated.

- Lacking any other information, the arithmetic mean will become best **predictor** for the variable under investigation.
- The deviations from the mean are the **unexplained** part or the **residuals** of the observations, i.e., $y_i = \bar{y} + \varepsilon_i$.



- Definition of total sum of squares: $TSS = \sum_{i=1}^n (Y_i - \bar{Y})^2$ or $TSS = \sum_{i=1}^n Y_i^2 - n \cdot \bar{Y}^2$.

- Why is the population variance estimated with $(n - 1)$ in the denominator, that is, by $s^2 = TSS/(n - 1)$:

Explanation 1: If we calculate the mean from the sample then there are only $n - 1$ "degrees of freedom" left because of the **zero sum property** of the mean.

Explanation 2: The mean is calculated by minimizing the *TSS*.

Thus the sample mean always fits the observed sample data better than any **unobserved but true** population expectation μ .

For the true expectation μ , the TSS would be slightly larger. That is why the sample TSS needs to be inflated by dividing it by a slight smaller value than n , that is, $n - 1$.

- **Standard deviation** measures the variation in **original units** rather than in squared units.

REVIEW: SKEWNESS

- Why does the distribution of the water consumption in the Concord dataset deviate from the normal distribution?

Reason: Fixed lower bound (negative water consumption impossible).

- **Skewness** and bounded/truncated distributions: For skewed distributions the notion of the center of the distribution (mean) becomes ambiguous and the **median** may be a better representation of the central tendency in the data.

- The **skewness** is defined by $skew(X) \equiv \frac{\sum_{i=1}^n (x_i - \bar{x})^3}{n \cdot s_X^3}$

- The normal distribution has a skewness of 0.

BOX-COX TRANSFORMATION

- This lecture focuses on the more general **Box-Cox** transformation (note 11 on page 28 in Hamilton) rather than the slightly simpler **power** transformation, which is discussed in Hamilton. For both transformations, the general interpretation of the parameter λ does not change.
- Causes for **extreme observations**: [a] skewed distributions, [b] measurement or recoding errors, [c] extreme but feasible events (perhaps not belonging to the population under investigation).
- The **power**-transformation presented in book and the **Box-Cox** transformation only work for **variables whose observations are all larger than zero**.
- The Box-Cox transformation is a generalization of the power transformation: $Y = \frac{X^\lambda - 1}{\lambda}$ and for $\lambda = 0$ we get $y = \ln(x)$.
- $\lambda > 1$ reduce negative skewness, whereas $\lambda < 1$ reduce positive skewness.
Remember: Positive skewness is very common for variables with a natural bound of zero.
- If power $\lambda < 0$ then all values are multiplied by a negative number to preserve the natural order of observations.

This explains the value λ in the denominator of the Box-Cox transformation

FOX Fig 4.1




Figure 4.1. The family of power transformations X' of X . The curve labeled p is the transformation $X^{(p)}$, that is, $(X^p - 1)/p$; $X^{(0)}$ is $\log_e X$.

- Note: R's function `car::powerTransform()` is performing several statistical tests whether a variable either needs to be transformed or whether a *log*-transformation is sufficient by using the likelihood ratio test (LR) principle:
 - The first LR tests the null hypotheses $H_0: \lambda^{optimal} = 0$. If we cannot reject the null hypothesis then we should tentatively work with a *log*-transformation to achieve normality/symmetry.
 - The second LR tests the null hypotheses $H_0: \lambda^{optimal} = 1$. If we cannot reject the null hypothesis then we should tentatively work with an untransformed variable because it is approximately symmetric.
 - The Wald confidence interval provides the 95% probability range within which the true population transformation parameter λ lies.

HANDLING TRANSFORMATIONS WITH NEGATIVE DATA VALUES

- After inspection of the variable's distribution one can overcome the problem of zero or negative data values by
 - adding a constant such as $\min(X) + \varepsilon$, where ε is a small positive number, or – say, 5% quantile – to the variable to make it solidly positive.
 - However, if ε is too small, which leads to positive but close to zero values, outliers may be introduced.
 - On the other hand, choosing ε too large, may make the transformation to normality ineffective.
- See the `?car::bcPower()` and Fox & Weissberg pp 161-162 for the **bcnPower** transformation family.
- A more informed way avoiding some of the problems by just adding a constant is to first transform the data by:

$$z(X, \gamma) = \frac{(X + \sqrt{X^2 + \gamma^2})}{2} \text{ with}$$

- The transformation $z(X, \gamma)$ is monotonic (i.e., if $x_1 < x_2$, then $z(x_1, \gamma) < z(x_2, \gamma)$)
 - For large positive X relative to γ ($X \gg \gamma$) the transformation is approximately linear with $z(X, \gamma) \approx X$.
 - If $\gamma = 0$ then $z(X, \gamma) = X$ for $X > 0$ and $z(X, \gamma) = 0$ for $X \leq 0$.
- Subsequently, once the γ -parameter is determined a standard Box-Cox transformation is applied to $z(X, \gamma)$.

LOESS SMOOTHER OF $Y \sim X$ RELATIONSHIPS

- Many of R's scatterplot functions not only show a linear regression fit through the data cloud but also show a locally smoothed loess-curve:
 - In essence, a sliding window moves over the value range of the variable X.
 - In each window a local regression line is estimated.
 - These local window regression lines are “splined” together into the smooth loess curve over the whole value range of X




FIGURE 9.17
MA-plot with curve obtained with loess.




FIGURE 9.16
Illustration of how loess estimates a curve. Showing 12 steps of the process.